

# Identifying Stable Kernels - top-down model

Nicholas Mc Guire <safety@osadl.org>

August 1, 2017



# Outline

- Context
- A few notes on statistics
- Data preparation
- The kernel as a whole
- Looking at subsystems
- Conclusions

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

## Outline

Context

Statistics

Subsystems

Conclusion

# Why GNU/Linux for Safety

- Satisfy Demands:
  - **multicore** there is no real alternative
  - **security** demands of C2C/C2X
  - **performance** demands in cognitive systems -> multi-core
  - functional requirements of e.g. autonomous systems (HAD)
- Satisfy Certification
  - standardization (notably requirements)
  - established concepts
  - breadth of deployment
  - development model (DLC data)

Using well-selected FLOSS for safety has clear technical advantages - developing the reusable generic arguments is a key element in SIL2LinuxMP.

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# Complex software and Safety

*"As far as practicable the design shall keep the safety-related part of the software simple"* [61508-3 Ed 7.4.2.6]

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

**Context**

Statistics

Subsystems

Conclusion

# Complex software and Safety

*"As far as practicable the design shall keep the safety-related part of the software simple" [61508-3 Ed 7.4.2.6]*

*"This Standard does not discourage the use of software in safety-critical systems. When designed and implemented correctly, software is often the first, and sometimes the best, hazard detection and prevention mechanism in the system."*  
[NASA NPR 8719.13B 1.2]

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# Complex software and Safety

*"As far as practicable the design shall keep the safety-related part of the software simple"* [61508-3 Ed 7.4.2.6]

*"This Standard does not discourage the use of software in safety-critical systems. When designed and implemented correctly, software is often the first, and sometimes the best, hazard detection and prevention mechanism in the system."*  
[NASA NPR 8719.13B 1.2]

**61508 Ed 2 essentially is guidance on complexity management at all levels of the DLC**

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# What changed in the safety business ?

- Performance demands
  - UP -> SMP/MP
  - 32bit -> 64bit
  - Large caches
  - Large memory footprint
- Non-determinism: caches, BTB, dynamic wait-states/retries
- Power-management: frequency-scaling, variable latency
- Complex OS - migration, pinning, per-core ops, physical concurrency
- Security - all the fun you can have in life in one problem
- Automotive industry discovered functional safety !

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# Re-use not re-write

- Rewriting complex software is not the solution - the problem is **not** the code but requirements and design.
- Violates security best practices (use common core libraries, use high level languages)
- Significantly changing pre-existing elements invalidates field data
  - -> risk of breaking the process
  - -> risk of more latent bugs (code level)
- Rewriting Linux would be a 60B USD waste
- 3S is not a fall back, it is the more adequate approach for complex systems if the elements selection is properly handled.

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion



# Level of complexity: Simple system calls

10000 runs - single call-tree

`sys_arch_prctl`

`sys_geteuid`

`Sys_getrlimit`

`Sys_lseek`

`Sys_rt_sigaction`

`Sys_sched_get_priority_max Sys_sched_get_priority_min`

`Sys_set_robust_list`

`Sys_brk`

`sys_getppid`

`sys_getuid`

`Sys_read`

`Sys_rt_sigprocmask`

Calls are simple:

- => Code Review doable
- => Testing doable with small test-samples
- => Field data may be quite relevant

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline


Context

Statistics

Subsystems

Conclusion

# Mid complexity system calls



- => Code review effort excessive
- => Need significantly more than 10k samples

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# Statistics

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Goal:

- Covering residual systematic faults
- Enabling architectural mitigations
- Managing/quantifying uncertainty -> recalls are costly !
- Address the limitations of highly-complex systems
- Provide key evidence for pre-existing elements
- Address business risks

"Assessment of non-compliant development", that is the demonstration of the adequacy of an element "as-is" and the process behind it, that was **not** initially intended for high-assurance/safety-related systems.

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# Some Basics

- Software exercised without change should show an exponential decrease of faults found
- Its a count -> poisson distribution
- But its highly overdispersed -> switch to negativ-binomial distribution
- Only view data-points -> borrowing strength over kernel versions.

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Clean approach:

The problem of valid/relevant metrics

- postulate your hypothesis
- design test metric
- sample  $\rightarrow$  statistic
- confirm/reject based on quantified criteria (e.g.  $\alpha$  level)
- Not realistic unfortunately because we do not yet know what we are looking for in many cases and complex systems rarely have simple cause-consequence models that are relevant.

The "clean" solution is the reference - but we know we must divert from this.

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Prediction on kernel - Data

What do we want to predict ?

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Prediction on kernel - Data

What do we want to predict ? **residual bugs in the kernel**

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion



# Prediction on kernel - Data

What do we want to predict ? **residual bugs in the kernel**

- Ideal input: time of bug-discovery

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Prediction on kernel - Data

What do we want to predict ? **residual bugs in the kernel**

- Ideal input: time of bug-discovery
- Real input:
  - time of bug-fix in -stable
  - -stable releases -> bug-fix only
  - Fixes: tag -> indication of bug source
  - Git data -> approximation of failure records
  - Uncertainty: trending over versions
  - Data-size issues: borrowing strenght.

Lets look at some examples

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Dirty approach:

- hypothesis -> plausible assumptions
- analyze available indirect data: e.g. bugs -> bug-fixes
- analyze available metrics -> data-base
- approximate/biased sample -> assess assumptions
- for relevant effects -> clean approach based on SRS
- verify process/methods/metrics by independent CA
- continuous monitoring to ensure validity of claims

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Dirty approach:

- hypothesis -> plausible assumptions
- analyze available indirect data: e.g. bugs -> bug-fixes
- analyze available metrics -> data-base
- approximate/biased sample -> assess assumptions
- for relevant effects -> clean approach based on SRS
- verify process/methods/metrics by independent CA
- continuous monitoring to ensure validity of claims
  - Risk: you might have to pull the plug because you violate your metrics

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Dirty approach:

- hypothesis -> plausible assumptions
- analyze available indirect data: e.g. bugs -> bug-fixes
- analyze available metrics -> data-base
- approximate/biased sample -> assess assumptions
- for relevant effects -> clean approach based on SRS
- verify process/methods/metrics by independent CA
- continuous monitoring to ensure validity of claims
  - Risk: you might have to pull the plug because you violate your metrics ... no need to wait until you kill someone

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

**Statistics**

Subsystems

Conclusion

# Limitations

- Independence violations: clustered findings/fixes
- Temporal uncertainties: time is **NOT** continuous
- zero-truncation and zero-inflation
- evolution of underlying process
- indirect data -> increased error terms
- small data sets -> borrowing strength
- Nonhomogenous population -> heteroscedasticity issues
- Data uncertainty: timestamp 1.1.1970, 14.2.2038 ?
- ...

Real data is not clean -> RIDM: manage/quantify uncertainty

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

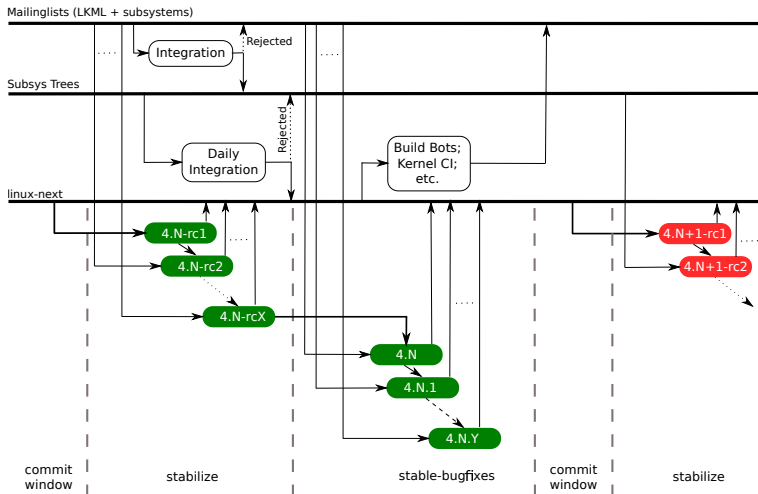
Context

Statistics

Subsystems

Conclusion

# Why we can do this - Linux DLC Stability



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# relating DLC to changes

## 4.4 life-cycle sequence

	files		lines		commits	lines added per commit
	changed	add	del			
rc1	9981	697599	460116	12226		57.0
rc2	334	4149	5497	386		10.7
rc3	245	3346	2342	277		12.0
rc4	363	4968	1672	331		15.0
rc5	256	1766	1304	260		6.7
rc6	263	2272	1236	309		7.3
rc7	91	977	429	109		8.9
rc8	73	709	448	82		8.6
4.4	88	518	280	102		5.0
4.4.1	80	644	280	120		5.3
4.4.2	112	1680	568	136		13.3
4.4.3	140	1307	585	343		3.8

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context


Statistics

Subsystems

Conclusion



# DLC process stability over Versions



Source: <http://neuling.org/linux-next>

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline


Context

Statistics

Subsystems

Conclusion

# Assessing suitability of Fixes: tag



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline



Context

Statistics

Subsystems

Conclusion

# Assessing suitability of Fixes: tag



Note that the Fixes: tag is not yet that well established (more or this later)

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# Analyze strength of coupling

A linear model could be a problem - the use of "Fixes:" is steadily increasing - but the increase within a single SUBLEVEL does not seem to be critical.

```
Call: lm(formula = bugs ~ fixes, data = d)

Residuals:
    Min       1Q   Median       3Q      Max
-59.627 -22.001  -4.902  16.287  76.149

Coefficients:
            Estimate Std. Error t value Pr(>|t|) Signif
(Intercept)   6.7077     21.1989   0.316   0.757
fixes         2.6014      0.4158   6.257 4.21e-05 0.001

Residual standard error: 38.92 on 12 degrees of freedom
Multiple R-squared:  0.7654,    Adjusted R-squared:  0.7458
F-statistic: 39.15 on 1 and 12 DF,  p-value: 4.21e-05
```

Conclusion: using the Fixes: tags should be a fairly robust representative of the full population (of bugs)

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline


Context

Statistics

Subsystems

Conclusion

# Bug introduction time (Fixes: tag)



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

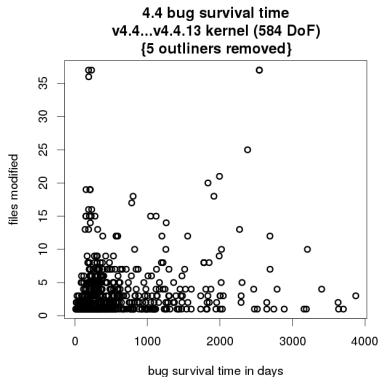

Context

Statistics

Subsystems

Conclusion

# Bug introduction time (Fixes: tag)



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline


Context

Statistics

Subsystems

Conclusion

# Distribution of bug survival in v4.4...v4.4.13



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline



Context

Statistics

Subsystems

Conclusion

# Prediction on kernel v4.4



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context





Statistics

Subsystems

Conclusion



# Prediction on kernel v4.4



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline





Context

Statistics

Subsystems

Conclusion

# Prediction on kernel v4.4 cont



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline





Context

Statistics

Subsystems

Conclusion

# Prediction on kernel v4.4 cont



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline


Context

Statistics

Subsystems

Conclusion

# Prediction on kernel v4.4 current



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline


Context

Statistics

Subsystems

Conclusion

# LTS non-LTS coupling



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline





Context

Statistics

Subsystems

Conclusion

# Prediction on kernel v4.9 cont



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline



Context

Statistics

Subsystems

Conclusion

# Prediction on kernel v4.9 cont



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline


Context

Statistics

Subsystems

Conclusion

# Prediction on kernel v4.9 current



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

[Outline](#)

[Context](#)





**[Statistics](#)**

[Subsystems](#)

[Conclusion](#)



# ...compare v3.18 ... v4.9



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# complete kernel - estimated residual fixes

ver	bugs/lv	bugs/hour	time-interval
3.18	1474	2011	3.16a
4.1	5154	6050	2.62a
4.4	1312	1468	1.81a
4.9	2496	2606	0.56a

Note these are simulation result snapshots (no error estimation yet)

We only can predict fixes (because that is what we measured) and we can only predict about as far into the future as we have robust data from the past - so widely varying prediction intervals.

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# estimated residual fixes densities

ver	files	fix/file	LoC	fix/LoC	density
3.18	46269	0.031 - 0.045	12711638	$1.1^{-4}$ — $1.6^{-4}$	1/7400
4.1	47661	0.108 - 0.127	13023108	$4.0^{-4}$ — $4.6^{-4}$	1/2300
4.4	50277	0.027 - 0.031	14023274	$9.3^{-5}$ — $1.0^{-4}$	1/10100
4.9	54069	0.046 - 0.048	14919875	$1.6^{-4}$ — $1.7^{-4}$	1/5600

Note these are simulation result snapshots (no error estimation yet) and not normalized ! - sorry did not realize that until I put the slides together

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# kernel DLC Trend v3.2 - v4.4(9)

Trending bugs-fixed over sublevel for -stable kernels

ver	slope	p-value	DoF	AIC
3.2	0.005910	0.06	83	904.76
3.4	0.0001224	<b>0.958</b>	112	1117
3.10	-0.004909	0.0166	103	1006.9
3.12	-0.002298	0.451	69	750.78
3.14	-0.014073	$6.26e - 07$	78	770.44
3.18	$-4.615e - 05$	0.00161	60	647.83
4.1	$7.656e - 06$	0.724	41	470.13
4.4	-0.014925	$1.15e - 05$	76	767.56
4.9	-0.005012	<b>0.585</b>	37	397.98

3.16 reappeared as stable at 3.16.35 but is not considered here as there is no adequate data for 3.16.8...3.16.35.

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline


Context

Statistics

Subsystems

Conclusion

# Parameter stability of prediction



**Identifying  
Stable  
Kernels -  
top-down  
model**

**Nicholas Mc  
Guire**  
<safety@osadl.

Outline





Context

**Statistics**

Subsystems

Conclusion

# Parameter stability of prediction



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion

# Stratifying

- identify high-risk subsystems over trending
- identify high-level dependencies/layering
- subset assessment based on specific configuration
- per patch statistics for subset/subsystem
- Testing:
  - temporal capabilities/interference
  - extreme-value distributions for conservative prediction
  - temporal monitoring at runtime -> SAC

Its not really clean any more as it is based on snooping the data technically a no-go: Mitigation -> assessment by independent authority in our case TueV Rheinland will review the statistical models and if they are valid.

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

**Subsystems**

Conclusion

# Subsystem Risk example

- /kernel vs full kernel
- /include vs /fs
- How fast does a subsystem stabilize -> identify risk
- How much does the full kernel "say" about a subsystem  
-> borrowing strength

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context





Statistics

**Subsystems**

Conclusion



# /kernel 3.18, 4.1, 4.4, 4.9 raw data



**Identifying  
Stable  
Kernels -  
top-down  
model**

**Nicholas Mc  
Guire**  
<safety@osadl.

**Outline**





**Context**

**Statistics**

**Subsystems**

**Conclusion**

# /fs 3.18, 4.1, 4.4, 4.9 raw data



**Identifying  
Stable  
Kernels -  
top-down  
model**

**Nicholas Mc  
Guire**  
<safety@osadl.

[Outline](#)





[Context](#)

[Statistics](#)

[Subsystems](#)

[Conclusion](#)

# full kernel vs /kernel 1 modeling



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline





Context

Statistics

Subsystems

Conclusion

# full kernel vs ./kernel cont.



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline





Context

Statistics

Subsystems

Conclusion

# full kernel vs ./kernel cont.



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline





Context

Statistics

Subsystems

Conclusion

# /include vs /fs 1 modeling



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline





Context

Statistics

Subsystems

Conclusion

# /include vs /fs cont.



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline





Context

Statistics

Subsystems

Conclusion

# /include vs /fs cont.



Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

Context

Statistics

Subsystems

Conclusion



# Conclusion

**Identifying  
Stable  
Kernels -  
top-down  
model**

**Nicholas Mc  
Guire  
<safety@osadl.**

Outline

Context

Statistics

Subsystems

**Conclusion**

# Minimize the config i7\_min\_config\_4.4.42-stable

Subsys	files	blank	comment	code	full	%used
arch	480	16725	20728	76208	2040204	3.74
block	52	3973	5619	16614	24753	67.12
crypto	32	1950	1618	8626	76265	11.31
drivers	357	32181	53023	128168	8587655	1.49
fs	138	14157	23640	80227	827737	9.69
include	1196	35986	64271	163952	449811	36.45
init	8	354	391	1846	2712	68.07
kernel	140	15662	28181	64968	161178	40.31
lib	97	2932	6816	16522	81891	20.18
mm	55	7521	15428	34409	70830	48.58
net	151	21472	17714	103505	650973	15.90
security	3	230	612	1127	50929	2.21

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

[Outline](#)

[Context](#)

[Statistics](#)

[Subsystems](#)

[Conclusion](#)

# Conclusions

- Relatively simple models allow identifying risks
- Trends of models **and** parameters are needed
- Subsystem models allow identifying less stable subsystems
- Trending over development allows detecting things going bad
- Comparison of kernel versions by objective data helps select the most reliable kernel - it is **not** though a single criteria that can give you an answer
- Minimizing your kernel is one simple elimination of residual bugs
- You need something stable ?: v3.18 or v4.4 would be my best guess at this point.

Statistic models may be a suitable extension for selecting kernels and other FLOSS elements to minimize risk. And clearly

- **The latest need not be the greatest**

Identifying  
Stable  
Kernels -  
top-down  
model

Nicholas Mc  
Guire  
<safety@osadl.

Outline

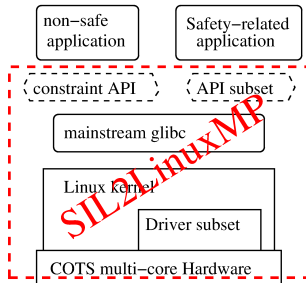
Context

Statistics

Subsystems

Conclusion

# The Goal



Thanks !